

# Amazon Fine Food Reviews Analysis

## [1]. Reading Data

### [1.1] Loading the data

```
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os

# Run this cell to mount your Google Drive.
from google.colab import drive
drive.mount('/content/drive')
```

➞ Go to this URL in a browser: [https://accounts.google.com/o/oauth2/auth?client\\_id=9473189](https://accounts.google.com/o/oauth2/auth?client_id=9473189)

Enter your authorization code:

.....

Mounted at /content/drive

```
!ls /content/drive/My\ Drive/Colab\ Notebooks
```

```

database.sqlite  Sample.ipynb      Untitled1.ipynb  Untitled4.ipynb
NB.ipynb        SVM.ipynb          Untitled2.ipynb
Reviews.csv     Untitled0.ipynb   Untitled3.ipynb

```

```
data=pd.read_csv('/content/drive/My Drive/Colab Notebooks/Reviews.csv')
```

```
data.head()
```

```

Id      ProductId      UserId  ProfileName  HelpfulnessNumerator  HelpfulnessDenominator

```

```
0    1    B001E4KFG0    A3SGXH7AUHU8GW    delmartian    1
```

```
1    2    B00813GRG4    A1D87F6ZCVE5NK    dll pa    0
```

```
conn=sqlite3.connect('/content/drive/My Drive/Colab Notebooks/database.sqlite')
```

```
filter_data=pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 50000""",conn)
```

```

def partition (x):
    if x<3:
        return 0
    return 1

```

```

actualscore = filter_data['Score']
positivenegative = actualscore.map(partition)
filter_data['Score']= positivenegative
print('Number of data points in our data',filter_data.shape)
filter_data.head(5)

```

```


```

Number of data points in our data (50000, 10)

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator
0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian		1

```
display = pd.read_sql_query("""
SELECT UserId, ProductId, ProfileName, Time, Score, Text, COUNT(*)
FROM Reviews
GROUP BY UserId
HAVING COUNT(*)>1
""", conn)
```

```
print(display.shape)
display.head()
```

(80668, 7)

	UserId	ProductId	ProfileName	Time	Score	Text
0	#oc-R115TNMSPFT9I7	B007Y59HVM	Breyton	1331510400	2	Overall its jus
1	#oc-R11D9D7SHXIJB9	B005HG9ET0	Louis E. Emory "hoppy"	1342396800	5	My wife ha
2	#oc-R11DNU2NBKQ23Z	B007Y59HVM	Kim Cieszykowski	1348531200	1	This coffee is ho
	"					"

```
display[display['UserId']=='AZY10LLTJ71NX']
```

	UserId	ProductId	ProfileName	Time	Score	Text
80638	AZY10LLTJ71NX	B006P7E5ZI	undertheshrine "undertheshrine"	1334707200	5	I was rec

```
display['COUNT(*)'].sum()
```

393063

## ▾ [2] Exploratory Data Analysis

## ▼ [2.1] Data Cleaning: Deduplication

```
display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND UserId="AR5J8UI46CURR"
ORDER BY ProductID
""", conn)
print(display.shape)
display.head()
```

↗ (5, 10)

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	Helpfulness
0	78445	B000HDL1RQ	AR5J8UI46CURR	Geetha Krishnan	2	
1	138317	B000HDOPYC	AR5J8UI46CURR	Geetha Krishnan	2	

Geetha

```
#Sorting data according to ProductId in ascending order
sorted_data=filter_data.sort_values('ProductId', axis=0, ascending=True, inplace=False, kind='quicksort')
```

```
#Deduplication of entries
final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"}, keep='first', inplace=False)
final.shape
```

↗ (46072, 10)

```
#Checking to see how much % of data still remains
(final['Id'].size*1.0)/(filter_data['Id'].size*1.0)*100
```

↗ 92.144

```
display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND Id=44737 OR Id=64422
ORDER BY ProductID
""", conn)

display.head()
```

↗

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Text
--	----	-----------	--------	-------------	----------------------	------------------------	-------	------

```
final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
```

0	04422	B000IMIDROQ	A101DR00JMC1F	Stephens	5	5	5	
---	-------	-------------	---------------	----------	---	---	---	--

```
#Before starting the next phase of preprocessing lets see the number of entries left
print(final.shape)
```

```
#How many positive and negative reviews are present in our dataset?
final['Score'].value_counts()
```

```
(46071, 10)
1      38479
0       7592
Name: Score, dtype: int64
```

## ▼ [3] Preprocessing

### ▼ [3.1]. Preprocessing Review Text

```
# printing some random reviews
sent_0 = final['Text'].values[0]
print(sent_0)
print("="*50)

sent_1000 = final['Text'].values[1000]
print(sent_1000)
print("="*50)
```

```
sent_1500 = final['Text'].values[1500]
print(sent_1500)
print("="*50)
```

```
sent_4900 = final['Text'].values[4900]
print(sent_4900)
print("="*50)
```

```
My dogs loves this chicken but its a product from China, so we wont be buying it anymore
=====
this is yummy, easy and unusual. it makes a quick, delicious pie, crisp or cobbler. home
=====
Great flavor, low in calories, high in nutrients, high in protein! Usually protein powde
=====
For those of you wanting a high-quality, yet affordable green tea, you should definitely
=====
```

```
# remove urls from text python: https://stackoverflow.com/a/40823105/4084039
sent_0 = re.sub(r"http\S+", "", sent_0)
sent_1000 = re.sub(r"http\S+", "", sent_1000)
sent_1500 = re.sub(r"http\S+", "", sent_1500)
sent_4900 = re.sub(r"http\S+", "", sent_4900)

print(sent_0)
```

☞ My dogs loves this chicken but its a product from China, so we wont be buying it anymore

```
# https://stackoverflow.com/questions/16206380/python-beautifulsoup-how-to-remove-all-tags-from-an-e
from bs4 import BeautifulSoup

soup = BeautifulSoup(sent_0, 'lxml')
text = soup.get_text()
print(text)
print("=="*50)

soup = BeautifulSoup(sent_1000, 'lxml')
text = soup.get_text()
print(text)
print("=="*50)

soup = BeautifulSoup(sent_1500, 'lxml')
text = soup.get_text()
print(text)
print("=="*50)

soup = BeautifulSoup(sent_4900, 'lxml')
text = soup.get_text()
print(text)
```

☞ My dogs loves this chicken but its a product from China, so we wont be buying it anymore  
 =====  
 this is yummy, easy and unusual. it makes a quick, delicious pie, crisp or cobbler. home  
 =====  
 Great flavor, low in calories, high in nutrients, high in protein! Usually protein powde  
 =====  
 For those of you wanting a high-quality, yet affordable green tea, you should definitely

```
# https://stackoverflow.com/a/47091490/4084039
import re

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)
    phrase = re.sub(r"'\re", " are", phrase)
    phrase = re.sub(r"'\s", " is", phrase)
    phrase = re.sub(r"'\d", " would", phrase)
    phrase = re.sub(r"'\ll", " will", phrase)
    phrase = re.sub(r"'\t", " not", phrase)
    phrase = re.sub(r"'\ve", " have", phrase)
    phrase = re.sub(r"'\m", " am", phrase)
    return phrase

sent_1500 = decontracted(sent_1500)
print(sent_1500)
print("=="*50)
```

☞

Great flavor, low in calories, high in nutrients, high in protein! Usually protein powde  
 =====

```
#remove words with numbers python: https://stackoverflow.com/a/18082370/4084039
sent_0 = re.sub("\S*\d\S*", "", sent_0).strip()
print(sent_0)
```

☞ My dogs loves this chicken but its a product from China, so we wont be buying it anymore

```
# https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'not'
# <br /><br /> ==> after the above steps, we are getting "br br"
# we are including them into stop words list
# instead of <br /> if we have <br/> these tags would have revmoved in the 1st step

stopwords= set(['br', 'the', 'i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "yo",
  "you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'him', 'his', 'himse',
  'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself', 'they', 'them',
  'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', "that'll", 'thes',
  'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'had', 'having',
  'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as', 'until', 'w',
  'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through', 'during', '
  'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off', 'over', 'under',
  'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'any', 'both', 'e',
  'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 'too', 'very', \
  's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've", 'now', 'd', 'll',
  've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't", 'doesn', "do",
  "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'mightn', "mightn",
  "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't", 'wasn', "wasn't",
  'won', "won't", 'wouldn', "wouldn't"])
```

```
# Combining all the above stundents
from tqdm import tqdm
preprocessed_reviews = []
# tqdm is for printing the status bar
for sentence in tqdm(final['Text'].values):
    sentence = re.sub(r"http\S+", "", sentence)
    sentence = BeautifulSoup(sentence, 'lxml').get_text()
    sentence = decontracted(sentence)
    sentence = re.sub("\S*\d\S*", "", sentence).strip()
    sentence = re.sub('[^A-Za-z]+', ' ', sentence)
    # https://gist.github.com/sebleier/554280
    sentence = ' '.join(e.lower() for e in sentence.split() if e.lower() not in stopwords)
    preprocessed_reviews.append(sentence.strip())
```

☞ 100%|██████████| 46071/46071 [00:17<00:00, 2583.23it/s]

```
preprocessed_reviews[1500]
```

☞ or low calories high nutrients high protein usually protein powders high priced high calo

## ▸ [4] Featurization

## ▼ [4.1] BAG OF WORDS

```
#Bow
count_vect = CountVectorizer() #in scikit-learn
count_vect.fit(preprocessed_reviews)
print("some feature names ", count_vect.get_feature_names()[:10])
print('='*50)

final_counts = count_vect.transform(preprocessed_reviews)
print("the type of count vectorizer ",type(final_counts))
print("the shape of out text BOW vectorizer ",final_counts.get_shape())
print("the number of unique words ", final_counts.get_shape()[1])
```

```
↳ some feature names ['aa', 'aaa', 'aaaa', 'aaaaa', 'aaaaaaaaaaaa', 'aaaaaaaaaaaaaaaa', 'a
=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (46071, 39364)
the number of unique words 39364
```

## ▼ [4.3] TF-IDF

```
tf_idf_vect = TfidfVectorizer(ngram_range=(1,2), min_df=10)
tf_idf_vect.fit(preprocessed_reviews)
print("some sample features(unique words in the corpus)",tf_idf_vect.get_feature_names()[0:10])
print('='*50)

final_tf_idf = tf_idf_vect.transform(preprocessed_reviews)
print("the type of count vectorizer ",type(final_tf_idf))
print("the shape of out text TFIDF vectorizer ",final_tf_idf.get_shape())
print("the number of unique words including both unigrams and bigrams ", final_tf_idf.get_shape()[1])
```

```
↳ some sample features(unique words in the corpus) ['ability', 'able', 'able buy', 'able c
=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (46071, 27311)
the number of unique words including both unigrams and bigrams 27311
```

## ▼ [4.4] Word2Vec

```
# Train your own Word2Vec model using your own text corpus
i=0
list_of_sentence=[]
for sentence in preprocessed_reviews:
    list_of_sentence.append(sentence.split())
```

```
# Using Google News Word2Vectors
```

```
# in this project we are using a pretrained model by google
# its 3.3G file, once you load this into your memory
# it occupies ~9Gb, so please do this step only if you have >12G of ram
# we will provide a pickle file wich contains a dict ,
# and it contains all our courpus words as keys and model[word] as values
# To use this code-snippet, download "GoogleNews-vectors-negative300.bin"
# from https://drive.google.com/file/d/0B7XkCwpI5KDYNlNUTTlSS21pQmM/edit
# it's 1.9GB in size.
```



```
# http://kavita-ganesan.com/gensim-word2vec-tutorial-starter-code/#.W17SRFAzZPY
# you can comment this whole cell
# or change these variable according to your need

is_your_ram_gt_16g=False
want_to_use_google_w2v = False
want_to_train_w2v = True

if want_to_train_w2v:
    # min_count = 5 considers only words that occurred atleast 5 times
    w2v_model=Word2Vec(list_of_sentence,min_count=5,size=50, workers=4)
    print(w2v_model.wv.most_similar('great'))
    print('='*50)
    print(w2v_model.wv.most_similar('worst'))

elif want_to_use_google_w2v and is_your_ram_gt_16g:
    if os.path.isfile('GoogleNews-vectors-negative300.bin'):
        w2v_model=KeyedVectors.load_word2vec_format('GoogleNews-vectors-negative300.bin', binary=True)
        print(w2v_model.wv.most_similar('great'))
        print(w2v_model.wv.most_similar('worst'))
    else:
        print("you don't have gogole's word2vec file, keep want_to_train_w2v = True, to train your o

☞ [('awesome', 0.816619873046875), ('fantastic', 0.804206907749176), ('terrific', 0.800173
=====
[('best', 0.7121520042419434), ('nastiest', 0.706802487373352), ('greatest', 0.704696536

w2v_words = list(w2v_model.wv.vocab)
print("number of words that occurred minimum 5 times ",len(w2v_words))
print("sample words ", w2v_words[0:50])

☞ number of words that occurred minimum 5 times 12798
sample words ['dogs', 'loves', 'chicken', 'product', 'china', 'wont', 'buying', 'anymor
```

## ▼ [4.4.1] Converting text into vectors using Avg W2V, TFIDF-W2V

### ▼ [4.4.1.1] Avg W2v

```
# average Word2Vec
# compute average word2vec for each review.
sent_vectors = []; # the avg-w2v for each sentence/review is stored in this list
for sent in tqdm(list_of_sentence): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might need to change this t
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    sent_vectors.append(sent_vec)
print(len(sent_vectors))
print(len(sent_vectors[0]))
```

```

↳ 100%|██████████| 46071/46071 [01:30<00:00, 508.05it/s]46071
50

```

#### ▼ [4.4.1.2] TFIDF weighted W2v

```

# S = ["abc def pqr", "def def def abc", "pqr pqr def"]
model = TfidfVectorizer()
tf_idf_matrix = model.fit_transform(preprocessed_reviews)
# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))

# TF-IDF weighted Word2Vec
tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in tqdm(list_of_sentence): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
            # tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole corpus
            # sent.count(word) = tf value of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_sent_vectors.append(sent_vec)
    row += 1

```

```

↳ 100%|██████████| 46071/46071 [16:34<00:00, 43.86it/s]

```

## ▼ [5] SVM

### ▼ [5.1] Linear SVM

#### ▼ [5.1.1] Applying Linear SVM on BOW

```

from sklearn.model_selection import train_test_split
from sklearn import preprocessing

#Breaking into Train and test
X_train, X_test, y_train, y_test = train_test_split(final_counts,final['Score'].values,test_size=0.2)
X_train, X_cv, y_train, y_cv = train_test_split(X_train, y_train, test_size=0.2)

#Normalize Data

```

```
X_train = preprocessing.normalize(X_train)
print("Train Data Size: ",X_train.shape)
```

```
#Normalize Data
```

```
X_test = preprocessing.normalize(X_test)
print("Test Data Size: ",X_test.shape)
```

```
X_cv = preprocessing.normalize(X_cv)
print("CV Data Size :", X_cv.shape)
```

```
↳ Train Data Size: (29484, 39364)
   Test Data Size: (9215, 39364)
   CV Data Size : (7372, 39364)
```

```
from sklearn.linear_model import SGDClassifier
from sklearn.calibration import CalibratedClassifierCV
from sklearn.metrics import roc_auc_score
import math
```

```
alpha = [10000,1000,100,10,1,0.1,0.01,0.001,0.0001]
```

```
train_auc = []
cv_auc = []
```


```
for i in alpha:
    model = SGDClassifier(alpha=i, loss = "hinge")
    clf = CalibratedClassifierCV(model, cv=3)
    clf.fit(X_train,y_train)
    prob_cv = clf.predict_proba(X_cv)[: ,1]
    cv_auc.append(roc_auc_score(y_cv,prob_cv))
    prob_train = clf.predict_proba(X_train)[: ,1]
    train_auc.append(roc_auc_score(y_train,prob_train))
optimal_alpha= alpha[cv_auc.index(max(cv_auc))]
alpha=[math.log(x) for x in alpha]
```

```
#plot auc vs alpha
```

```
x = plt.subplot( )
x.plot(alpha, train_auc, label='AUC train')
x.plot(alpha, cv_auc, label='AUC CV')
plt.title('AUC vs hyperparameter')
plt.xlabel('alpha')
plt.ylabel('AUC')
x.legend()
plt.show()
```

```
print('optimal alpha for which auc is maximum : ',optimal_alpha)
```

```
↳
```



#Testing AUC on Test data

```
model = SGDCClassifier(alpha = optimal_alpha)
clf = CalibratedClassifierCV(model, cv=3)
clf.fit(X_train,y_train)
pred_test = clf.predict_proba(X_test)[: ,1]
fpr1, tpr1, thresholds1 = metrics.roc_curve(y_test, pred_test)
pred_train = clf.predict_proba(X_train)[: ,1]
fpr2,tpr2,thresholds2 = metrics.roc_curve(y_train,pred_train)
```

#plot ROC curve

```
x = plt.subplot( )
x.plot(fpr1, tpr1, label = 'Test ROC ,auc='+str(roc_auc_score(y_test,pred_test)))
x.plot(fpr2, tpr2, label='Train ROC ,auc='+str(roc_auc_score(y_train,pred_train)))
plt.title('ROC curve')
plt.xlabel('FPR')
plt.ylabel('TPR')
x.legend()
plt.show()
```


```
print("AUC on Test data is " +str(roc_auc_score(y_test,pred_test)))
print("AUC on Train data is " +str(roc_auc_score(y_train,pred_train)))
```

```
print("-----")
```

# Code for drawing seaborn heatmaps


```
class_names = ['negative','positive']
df_heatmap = pd.DataFrame(confusion_matrix(y_test, pred_test.round()), index=class_names, columns=c1)
fig = plt.figure( )
heatmap = sns.heatmap(df_heatmap, annot=True, fmt="d")
```





AUC on Test data is 0.9455923816450589

AUC on Train data is 0.9624459384636701



```
results=pd.DataFrame(columns=['Featuraization', 'Classifier', 'alpha', 'Train-AUC', 'Test-AUC' ])
new = ['BOW', 'SGDClassifier-hinge loss', 0.0001, 0.9624, 0.9455.]
results.loc[0] = new
```

```
from sklearn.linear_model import SGDClassifier
```

```
# top 10 positive features
features = count_vect.get_feature_names()
clf = SGDClassifier(alpha=0.0001)
clf.fit(X_train,y_train)
weight = clf.coef_
positive_index=np.argsort(weight)[::-1]
```

```
negative_index=np.argsort(weight)
```

```
print('Top 10 positive features :')
for i in list(positive_index[0][0:10]):
    print(features[i])
```

```
print("-----")
```

```
#top 10 negative features
print('Top 10 negative features :')
for i in list(negative_index[0][0:10]):
    print(features[i])
```



Top 10 positive features :

delicious  
best  
great  
nice  
amazing  
loves  
perfect  
wonderful  
excellent  
love

-----

Top 10 negative features :

worst  
disappointed  
horrible  
terrible  
disappointing  
awful  
return

## ▼ [5.1.2] Applying Linear SVM on TFIDF

```
X_train, X_test, y_train, y_test = train_test_split(final_tf_idf, final['Score'].values, test_size=0.2)
X_train, X_cv, y_train, y_cv = train_test_split(X_train, y_train, test_size=0.2)
```

```
#Normalize Data
```

```
X_train = preprocessing.normalize(X_train)
print("Train Data Size: ", X_train.shape)
```

```
#Normalize Data
```

```
X_test = preprocessing.normalize(X_test)
print("Test Data Size: ", X_test.shape)
```

```
#Normalize Data
```

```
X_cv = preprocessing.normalize(X_cv)
print("CV Data size:", X_cv.shape)
```

```
↳ Train Data Size: (29484, 27311)
   Test Data Size: (9215, 27311)
   CV Data size: (7372, 27311)
```

```
alpha = [10000, 1000, 100, 10, 1, 0.1, 0.01, 0.001, 0.0001]
```

```
train_auc = []
cv_auc = []
```

```
for i in alpha:
    model = SGDClassifier(alpha=i, loss = "hinge")
    clf = CalibratedClassifierCV(model, cv=3)
    clf.fit(X_train, y_train)
    prob_cv = clf.predict_proba(X_cv)[:, 1]
    cv_auc.append(roc_auc_score(y_cv, prob_cv))
    prob_train = clf.predict_proba(X_train)[:, 1]
    train_auc.append(roc_auc_score(y_train, prob_train))
optimal_alpha = alpha[cv_auc.index(max(cv_auc))]
alpha = [math.log(x) for x in alpha]
```


```
#plot auc vs alpha
x = plt.subplot( )
```

```

x.plot(alpha, train_auc, label='AUC train')
x.plot(alpha, cv_auc, label='AUC CV')
plt.title('AUC vs hyperparameter')
plt.xlabel('alpha')
plt.ylabel('AUC')
x.legend()
plt.show()

print('optimal alpha for which auc is maximum : ',optimal_alpha)

```



optimal alpha for which auc is maximum : 0.0001

```

#Testing AUC on Test data
model = SGDCClassifier(alpha = optimal_alpha)
clf = CalibratedClassifierCV(model, cv=3)
clf.fit(X_train,y_train)
pred_test = clf.predict_proba(X_test)[: ,1]
fpr1, tpr1, thresholds1 = metrics.roc_curve(y_test, pred_test)
pred_train = clf.predict_proba(X_train)[: ,1]
fpr2,tpr2,thresholds2 = metrics.roc_curve(y_train,pred_train)

#plot ROC curve
x = plt.subplot( )
x.plot(fpr1, tpr1, label = 'Test ROC ,auc='+str(roc_auc_score(y_test,pred_test)))
x.plot(fpr2, tpr2, label='Train ROC ,auc='+str(roc_auc_score(y_train,pred_train)))
plt.title('ROC curve')
plt.xlabel('FPR')
plt.ylabel('TPR')
x.legend()
plt.show()


print("AUC on Test data is " +str(roc_auc_score(y_test,pred_test)))
print("AUC on Train data is " +str(roc_auc_score(y_train,pred_train)))

print("-----")

# Code for drawing seaborn heatmaps
class_names = ['negative','positive']
df_heatmap = pd.DataFrame(confusion_matrix(y_test, pred_test.round()), index=class_names, columns=cl
fig = plt.figure( )
heatmap = sns.heatmap(df_heatmap, annot=True, fmt="d")


```





AUC on Test data is 0.9586475114116522

AUC on Train data is 0.9841205419840111



```
new = ['tf-idf', 'SGDClassifier-hinge loss', 0.0001, 0.9841, 0.9586].
results.loc[1] = new
```

```
# top 10 positive features
features = tf_idf_vect.get_feature_names()
clf = SGDClassifier(alpha=0.0001)
clf.fit(X_train,y_train)
weight = clf.coef_
positive_index=np.argsort(weight)[:,:-1]

negative_index=np.argsort(weight)

print('Top 10 positive features :')
for i in list(positive_index[0][0:10]):
    print(features[i])

print("-----")

#top 10 negative features
print('Top 10 negative features :')
for i in list(negative_index[0][0:10]):
    print(features[i])
```





Top 10 positive features :

great  
good  
best  
not disappointed  
delicious  
love  
loves  
nice  
amazing  
wonderful

-----

Top 10 negative features :

disappointed  
worst  
disappointing  
horrible  
not  
terrible  
not recommend  
not worth

### ▼ [5.1.3] Applying Linear SVM on AVG W2V

```
avg_vec_google = np.array(sent_vectors)
```

```
mask = ~np.any(np.isnan(avg_vec_google), axis=1)
# print(mask)
avg_vec_google_new = avg_vec_google[mask]
df_sample_new = final['Score'][mask]
print(avg_vec_google_new.shape)
print(df_sample_new.shape)
```

```
↳ (46071, 50)
   (46071,)
```

```
#Normalizing the data
avg_vec_norm = preprocessing.normalize(avg_vec_google_new)

#Not shuffling the data as we want it on time basis
X_train, X_test, y_train, y_test = train_test_split(avg_vec_norm, df_sample_new.values, test_size=0.2,
X_train, X_cv, y_train, y_cv = train_test_split(X_train, y_train, test_size=0.2)

#Normalize Data
X_train = preprocessing.normalize(X_train)
print("Train Data Size: ", X_train.shape)

#Normalize Data
X_test = preprocessing.normalize(X_test)
print("Test Data Size: ", X_test.shape)

#Normalize Data
X_cv = preprocessing.normalize(X_cv)
print("CV Data size:", X_cv.shape)
```

```
↳
```

Train Data Size: (29484, 50)

```
alpha = [10000,1000,100,10,1,0.1,0.01,0.001,0.0001]
```

```
train_auc = []
```


```
cv_auc = []
```

```
for i in alpha:
    model = SGDClassifier(alpha=i, loss = "hinge")
    clf = CalibratedClassifierCV(model, cv=3)
    clf.fit(X_train,y_train)
    prob_cv = clf.predict_proba(X_cv)[: ,1]
    cv_auc.append(roc_auc_score(y_cv,prob_cv))
    prob_train = clf.predict_proba(X_train)[: ,1]
    train_auc.append(roc_auc_score(y_train,prob_train))
optimal_alpha= alpha[cv_auc.index(max(cv_auc))]
alpha=[math.log(x) for x in alpha]
```

```
#plot auc vs alpha
```

```
x = plt.subplot( )
x.plot(alpha, train_auc, label='AUC train')
x.plot(alpha, cv_auc, label='AUC CV')
plt.title('AUC vs hyperparameter')
plt.xlabel('alpha')
plt.ylabel('AUC')
x.legend()
plt.show()
```

```
print('optimal alpha for which auc is maximum : ',optimal_alpha)
```



optimal alpha for which auc is maximum : 0.0001

```
#Testing AUC on Test data
```

```
model = SGDClassifier(alpha = optimal_alpha)
clf = CalibratedClassifierCV(model, cv=3)
clf.fit(X_train,y_train)
pred_test = clf.predict_proba(X_test)[: ,1]
fpr1, tpr1, thresholds1 = metrics.roc_curve(y_test, pred_test)
pred_train = clf.predict_proba(X_train)[: ,1]
fpr2,tpr2,thresholds2 = metrics.roc_curve(y_train,pred_train)
```

```
#plot ROC curve
```

```
x = plt.subplot( )
x.plot(fpr1, tpr1, label = 'Test ROC ,auc='+str(roc_auc_score(y_test,pred_test)))
x.plot(fpr2, tpr2, label='Train ROC ,auc='+str(roc_auc_score(y_train,pred_train)))
plt.title('ROC curve')
plt.xlabel('FPR')
```

```


plt.ylabel('TPR')
x.legend()
plt.show()

print("AUC on Test data is " +str(roc_auc_score(y_test,pred_test)))
print("AUC on Train data is " +str(roc_auc_score(y_train,pred_train)))


print("-----")

# Code for drawing seaborn heatmaps
class_names = ['negative','positive']
df_heatmap = pd.DataFrame(confusion_matrix(y_test, pred_test.round()), index=class_names, columns=cl
fig = plt.figure( )
heatmap = sns.heatmap(df_heatmap, annot=True, fmt="d")

```



AUC on Test data is 0.9050066517886014  
AUC on Train data is 0.903404171453736  
-----



```

new = ['AVG W2V','SGDClassifier-hinge loss',0.0001,0.9034,0.9050]
results.loc[2] = new

```

#### ▼ [5.1.4] Applying Linear SVM on TFIDF W2V

```

tfidf_w2v_vec_google = np.array(tfidf_sent_vectors)

tfidfw2v_vecs_norm = preprocessing.normalize(tfidf_w2v_vec_google)
X_train, X_test, y_train, y_test = train_test_split(tfidfw2v_vecs_norm, final['Score'].values, test_size=0.2)
X_train, X_cv, y_train, y_cv = train_test_split(X_train, y_train, test_size=0.2)

alpha = [10000,1000,100,10,1,0.1,0.01,0.001,0.0001]


train_auc = []
cv_auc = []

for i in alpha:
    model = SGDClassifier(alpha=i, loss = "hinge")
    clf = CalibratedClassifierCV(model, cv=3)
    clf.fit(X_train,y_train)
    prob_cv = clf.predict_proba(X_cv)[:,-1]
    cv_auc.append(roc_auc_score(y_cv,prob_cv))
    prob_train = clf.predict_proba(X_train)[:,-1]
    train_auc.append(roc_auc_score(y_train,prob_train))
optimal_alpha= alpha[cv_auc.index(max(cv_auc))]
alpha=[math.log(x) for x in alpha]

#plot auc vs alpha
x = plt.subplot( )
x.plot(alpha, train_auc, label='AUC train')
x.plot(alpha, cv_auc, label='AUC CV')
plt.title('AUC vs hyperparameter')
plt.xlabel('alpha')
plt.ylabel('AUC')
x.legend()
plt.show()

print('optimal alpha for which auc is maximum : ',optimal_alpha)

```



optimal alpha for which auc is maximum : 0.0001

```

#Testing AUC on Test data
model = SGDClassifier(alpha = optimal_alpha)
clf = CalibratedClassifierCV(model, cv=3)
clf.fit(X_train,y_train)
pred_test = clf.predict_proba(X_test)[:,-1]
fpr1, tpr1, thresholds1 = metrics.roc_curve(y_test, pred_test)
pred_train = clf.predict_proba(X_train)[:,-1]
fpr2, tpr2, thresholds2 = metrics.roc_curve(y_train, pred_train)

```

```


#plot ROC curve
x = plt.subplot( )
x.plot(fpr1, tpr1, label='Test ROC ,auc='+str(roc_auc_score(y_test,pred_test)))
x.plot(fpr2, tpr2, label='Train ROC ,auc='+str(roc_auc_score(y_train,pred_train)))
plt.title('ROC curve')
plt.xlabel('FPR')
plt.ylabel('TPR')
x.legend()
plt.show()

print("AUC on Test data is " +str(roc_auc_score(y_test,pred_test)))
print("AUC on Train data is " +str(roc_auc_score(y_train,pred_train)))


print("-----")

# Code for drawing seaborn heatmaps
class_names = ['negative','positive']
df_heatmap = pd.DataFrame(confusion_matrix(y_test, pred_test.round()), index=class_names, columns=c1)
fig = plt.figure( )
heatmap = sns.heatmap(df_heatmap, annot=True, fmt="d")

```



AUC on Test data is 0.8785304507674528  
AUC on Train data is 0.8801411814227382



```
new = ['tf-idf W2V', 'SGDClassifier-hinge loss', 0.0001, 0.8801, 0.8785].
results.loc[3] = new
```

## ▼ [5.2] RBF SVM

### ▼ [5.2.1] Applying RBF SVM on BOW

```
X=preprocessed_reviews
y=np.array(final['Score'])
X=X[:30000]
y=y[:30000]
```

```
#BoW
count_vect = CountVectorizer(min_df=10, max_features=500) #in scikit-learn
count_vect.fit(preprocessed_reviews)
print("some feature names ", count_vect.get_feature_names()[:10])
print('='*50)
```

```
final_counts = count_vect.transform(preprocessed_reviews)
print("the type of count vectorizer ", type(final_counts))
print("the shape of out text BOW vectorizer ", final_counts.get_shape())
print("the number of unique words ", final_counts.get_shape()[1])
```

```

↳ some feature names ['able', 'absolutely', 'acid', 'actually', 'add', 'added', 'aftertas
=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text BOW vectorizer (46071, 500)
the number of unique words 500

```

```
#Breaking into Train and test
```

```
X_train, X_test, y_train, y_test = train_test_split(final_counts, final['Score'].values, test_size=0.2)
X_train, X_cv, y_train, y_cv = train_test_split(X_train, y_train, test_size=0.2)
```

```
#Normalize Data
```

```
X_train = preprocessing.normalize(X_train)
print("Train Data Size: ", X_train.shape)
```

```
#Normalize Data
```

```
X_test = preprocessing.normalize(X_test)
print("Test Data Size: ", X_test.shape)
```

```
X_cv = preprocessing.normalize(X_cv)
print("CV Data Size : ", X_cv.shape)
```

```

↳ Train Data Size: (29484, 500)
Test Data Size: (9215, 500)
CV Data Size : (7372, 500)

```

```
from sklearn.svm import SVC
```

```
C = [10000, 1000, 100, 10, 1, 0.1, 0.01, 0.001, 0.0001]
```

```
train_auc = []
```

```
cv_auc = []
```

```

for i in C:
    model = SVC(C=i)
    clf = CalibratedClassifierCV(model, cv=3)
    clf.fit(X_train, y_train)
    prob_cv = clf.predict_proba(X_cv)[:, 1]
    cv_auc.append(roc_auc_score(y_cv, prob_cv))
    prob_train = clf.predict_proba(X_train)[:, 1]
    train_auc.append(roc_auc_score(y_train, prob_train))
optimal_C = C[cv_auc.index(max(cv_auc))]
C = [math.log(x) for x in C]

```

```
#plot auc vs alpha
```

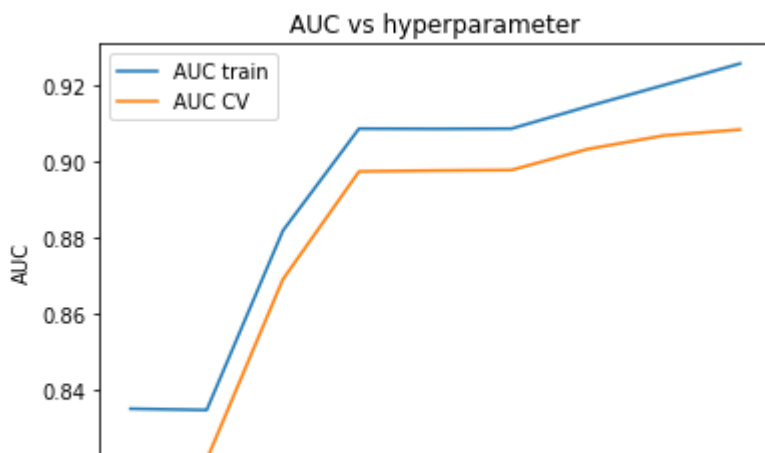
```

x = plt.subplot( )
x.plot(C, train_auc, label='AUC train')
x.plot(C, cv_auc, label='AUC CV')
plt.title('AUC vs hyperparameter')
plt.xlabel('C')
plt.ylabel('AUC')
x.legend()
plt.show()

```

```
print('optimal C for which auc is maximum : ', C)
```

```
↳
```



```
print('optimal C for which auc is maximum : ',optimal_C)
```

```
↳ optimal C for which auc is maximum : 10000
```

```
#Testing AUC on Test data
model =SVC(C = optimal_C)
clf = CalibratedClassifierCV(model, cv=3)
clf.fit(X_train,y_train)
pred_test = clf.predict_proba(X_test)[: ,1]
fpr1, tpr1, thresholds1 = metrics.roc_curve(y_test, pred_test)
pred_train = clf.predict_proba(X_train)[: ,1]
fpr2,tpr2,thresholds2 = metrics.roc_curve(y_train,pred_train)

#plot ROC curve
x = plt.subplot( )
x.plot(fpr1, tpr1, label = 'Test ROC ,auc='+str(roc_auc_score(y_test,pred_test)))
x.plot(fpr2, tpr2, label='Train ROC ,auc='+str(roc_auc_score(y_train,pred_train)))
plt.title('ROC curve')
plt.xlabel('FPR')
plt.ylabel('TPR')
x.legend()
plt.show()


print("AUC on Test data is " +str(roc_auc_score(y_test,pred_test)))
print("AUC on Train data is " +str(roc_auc_score(y_train,pred_train)))

print("-----")

# Code for drawing seaborn heatmaps
class_names = ['negative','positive']
df_heatmap = pd.DataFrame(confusion_matrix(y_test, pred_test.round()), index=class_names, columns=cl
fig = plt.figure( )
heatmap = sns.heatmap(df_heatmap, annot=True, fmt="d")
```


```
↳
```





AUC on Test data is 0.9100668107746447

AUC on Train data is 0.9256879098972174



```
results_2=pd.DataFrame(columns=['Featuraization', 'Classifier','C' , 'Train-AUC' , 'Test-AUC' ])
```

```
new = ["Bow", "RBF", 10000, 0.9256, 0.9100].
results_2.loc[0] = new
```

## ▼ [5.2.2] Applying RBF SVM on TFIDF

```
tf_idf_vect = TfidfVectorizer(ngram_range=(1,2), min_df=10,max_features=500)
tf_idf_vect.fit(preprocessed_reviews)
print("some sample features(unique words in the corpus)",tf_idf_vect.get_feature_names()[0:10])
print('='*50)

final_tf_idf = tf_idf_vect.transform(preprocessed_reviews)
print("the type of count vectorizer ",type(final_tf_idf))
print("the shape of out text TFIDF vectorizer ",final_tf_idf.get_shape())
print("the number of unique words including both unigrams and bigrams ", final_tf_idf.get_shape()[1])
```

```
↳ some sample features(unique words in the corpus) ['able', 'absolutely', 'acid', 'actual']
=====
the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of out text TFIDF vectorizer (46071, 500)
the number of unique words including both unigrams and bigrams 500
```

```
#Breaking into Train and test
```

```
X_train, X_test, y_train, y_test = train_test_split(final_counts,final['Score'].values,test_size=0.2)
X_train, X_cv, y_train, y_cv = train_test_split(X_train, y_train, test_size=0.2)
```

```
#Normalize Data
```

```
X_train = preprocessing.normalize(X_train)
print("Train Data Size: ",X_train.shape)
```

```
#Normalize Data
```

```
X_test = preprocessing.normalize(X_test)
print("Test Data Size: ",X_test.shape)
```

```
X_cv = preprocessing.normalize(X_cv)
print("CV Data Size :", X_cv.shape)
```

```
↳ Train Data Size: (29484, 500)
   Test Data Size: (9215, 500)
   CV Data Size : (7372, 500)
```

```
from sklearn.svm import SVC
```

```
C = [10000,1000,100,10,1,0.1,0.01,0.001,0.0001]
```

```
train_auc = []
cv_auc = []
```


```
for i in C:
    model = SVC(C=i)
    clf = CalibratedClassifierCV(model, cv=3)
    clf.fit(X_train,y_train)
    prob_cv = clf.predict_proba(X_cv)[: ,1]
    cv_auc.append(roc_auc_score(y_cv,prob_cv))
    prob_train = clf.predict_proba(X_train)[: ,1]
    train_auc.append(roc_auc_score(y_train,prob_train))
optimal_C= C[cv_auc.index(max(cv_auc))]
C=[math.log(x) for x in C]
```

```
#plot auc vs alpha
```

```
x = plt.subplot( )
x.plot(C, train_auc, label='AUC train')
x.plot(C, cv_auc, label='AUC CV')
plt.title('AUC vs hyperparameter')
plt.xlabel('C')
plt.ylabel('AUC')
x.legend()
plt.show()
```

```
print('optimal C for which auc is maximum : ',optimal_C)
```

```
↳
```



```
#Testing AUC on Test data
model =SVC(C = optimal_C)
clf = CalibratedClassifierCV(model, cv=3)
clf.fit(X_train,y_train)
pred_test = clf.predict_proba(X_test)[: ,1]
fpr1, tpr1, thresholds1 = metrics.roc_curve(y_test, pred_test)
pred_train = clf.predict_proba(X_train)[: ,1]
fpr2,tpr2,thresholds2 = metrics.roc_curve(y_train,pred_train)


#plot ROC curve
x = plt.subplot( )
x.plot(fpr1, tpr1, label = 'Test ROC ,auc='+str(roc_auc_score(y_test,pred_test)))
x.plot(fpr2, tpr2, label='Train ROC ,auc='+str(roc_auc_score(y_train,pred_train)))
plt.title('ROC curve')
plt.xlabel('FPR')
plt.ylabel('TPR')
x.legend()
plt.show()

print("AUC on Test data is " +str(roc_auc_score(y_test,pred_test)))
print("AUC on Train data is " +str(roc_auc_score(y_train,pred_train)))

print("-----")


# Code for drawing seaborn heatmaps
class_names = ['negative','positive']
df_heatmap = pd.DataFrame(confusion_matrix(y_test, pred_test.round()), index=class_names, columns=cl
fig = plt.figure( )
heatmap = sns.heatmap(df_heatmap, annot=True, fmt="d")
```





AUC on Test data is 0.9078005638360082

AUC on Train data is 0.9253188130489396



```
new = ["tf-idf", "RBF", 10000, 0.9253, 0.9078].
results_2.loc[1] = new
```

### ▼ [5.2.3] Applying RBF SVM on AVG W2V

```
# Train your own Word2Vec model using your own text corpus
```

```
i=0
```

```
list_of_sentence=[]
```

```
for sentence in X:
```

```
    list_of_sentence.append(sentence.split())
```

```
# min_count = 5 considers only words that occurred at least 5 times
```

```
w2v_model=Word2Vec(list_of_sentence,min_count=5,size=50,workers=4)
```

```
w2v_words = list(w2v_model.wv.vocab)
```

```
# average Word2Vec
```

```
# compute average word2vec for each review.
```

```
sent_vectors = []; # the avg-w2v for each sentence/review is stored in this list
```

```
for sent in tqdm(list_of_sentence): # for each review/sentence
```

```
    sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might need to change this t
```

```
    cnt_words = 0; # num of words with a valid vector in the sentence/review
```

```
    for word in sent: # for each word in a review/sentence
```

```
        if word in w2v_words:
```

```
            vec = w2v_model.wv[word]
```

```
            sent_vec += vec
```

```
            cnt_words += 1
```

```
    if cnt_words != 0:
```

```

        sent_vec /= cnt_words
    sent_vectors.append(sent_vec)
print(len(sent_vectors))
print(len(sent_vectors[0]))

```

```

↳ 100%|██████████| 30000/30000 [00:53<00:00, 555.74it/s]30000
50

```

```
avg_vec_google = np.array(sent_vectors)
```

```

mask = ~np.any(np.isnan(avg_vec_google), axis=1)
# print(mask)
avg_vec_google_new = avg_vec_google[mask]
df_sample_new = y[mask]
print(avg_vec_google_new.shape)
print(df_sample_new.shape)

```

```

↳ (30000, 50)
(30000,)

```

```

from sklearn import preprocessing
from sklearn.model_selection import train_test_split

```

```

#Normalizing the data
avg_vec_norm = preprocessing.normalize(avg_vec_google_new)

```

```

#Not shuffling the data as we want it on time basis
X_train, X_test, y_train, y_test = train_test_split(avg_vec_norm, df_sample_new, test_size=0.2, shuffle=False)
X_train, X_cv, y_train, y_cv = train_test_split(X_train, y_train, test_size=0.2)

```

```

#Normalize Data
X_train = preprocessing.normalize(X_train)
print("Train Data Size: ", X_train.shape)

```

```

#Normalize Data
X_test = preprocessing.normalize(X_test)
print("Test Data Size: ", X_test.shape)

```

```

#Normalize Data
X_cv = preprocessing.normalize(X_cv)
print("CV Data size:", X_cv.shape)

```

```

↳ Train Data Size: (19200, 50)
Test Data Size: (6000, 50)
CV Data size: (4800, 50)

```

```

from sklearn.svm import SVC
from sklearn.calibration import CalibratedClassifierCV
from sklearn.metrics import roc_auc_score
import math

```

```
C = [10000, 1000, 100, 10, 1, 0.1, 0.01, 0.001, 0.0001]
```

```

train_auc = []
cv_auc = []

```

```

for i in C:
    model = SVC(C=i)
    clf = CalibratedClassifierCV(model, cv=3)

```

```

clf.fit(X_train,y_train)
prob_cv = clf.predict_proba(X_cv)[: ,1]
cv_auc.append(roc_auc_score(y_cv,prob_cv))
prob_train = clf.predict_proba(X_train)[: ,1]
train_auc.append(roc_auc_score(y_train,prob_train))
optimal_C= C[cv_auc.index(max(cv_auc))]
C=[math.log(x) for x in C]

```

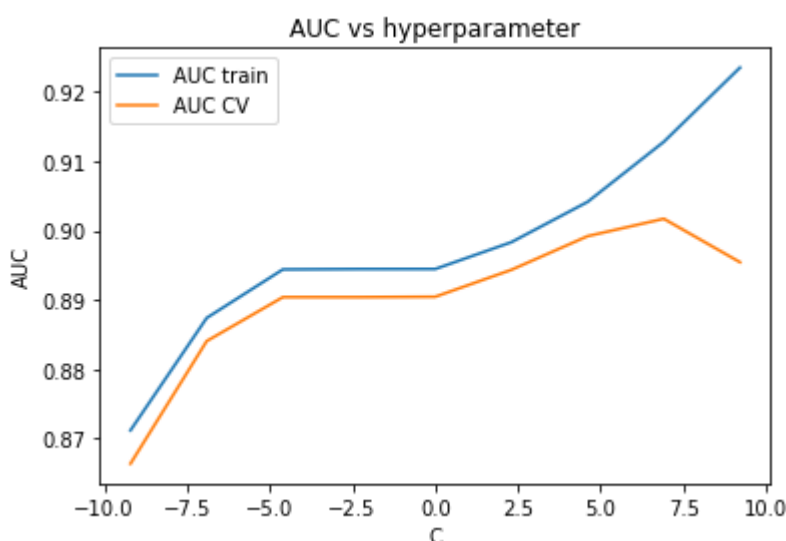
```
#plot auc vs alpha
```

```

x = plt.subplot( )
x.plot(C, train_auc, label='AUC train')
x.plot(C, cv_auc, label='AUC CV')
plt.title('AUC vs hyperparameter')
plt.xlabel('C')
plt.ylabel('AUC')
x.legend()
plt.show()

```

```
print('optimal C for which auc is maximum : ',optimal_C)
```



```
optimal C for which auc is maximum : 1000
```

```
from sklearn.metrics import confusion_matrix
```

```
#Testing AUC on Test data
```

```

model =SVC(C = optimal_C)
clf = CalibratedClassifierCV(model, cv=3)
clf.fit(X_train,y_train)
pred_test = clf.predict_proba(X_test)[: ,1]
fpr1, tpr1, thresholds1 = metrics.roc_curve(y_test, pred_test)
pred_train = clf.predict_proba(X_train)[: ,1]
fpr2,tpr2,thresholds2 = metrics.roc_curve(y_train,pred_train)

```

```
#plot ROC curve
```

```

x = plt.subplot( )
x.plot(fpr1, tpr1, label = 'Test ROC ,auc='+str(roc_auc_score(y_test,pred_test)))
x.plot(fpr2, tpr2, label='Train ROC ,auc='+str(roc_auc_score(y_train,pred_train)))
plt.title('ROC curve')
plt.xlabel('FPR')
plt.ylabel('TPR')
x.legend()
plt.show()

```


```

print("AUC on Test data is " +str(roc_auc_score(y_test,pred_test)))
print("AUC on Train data is " +str(roc_auc_score(y_train,pred_train)))


```

```
print("-----")
```

```
# Code for drawing seaborn heatmaps
class_names = ['negative', 'positive']
df_heatmap = pd.DataFrame(confusion_matrix(y_test, pred_test.round()), index=class_names, columns=cl
fig = plt.figure( )
heatmap = sns.heatmap(df_heatmap, annot=True, fmt="d")
```



AUC on Test data is 0.9058774651930146  
AUC on Train data is 0.9128031475958928



```
new = ["AVG W2V", "RBF", 1000, 0.9128, 0.9058.]
results_2.loc[2] = new
```

## ▼ [5.2.4] Applying RBF SVM on TFIDF W2V

```
# S = ["abc def pqr", "def def def abc", "pqr pqr def"]
model = TfidfVectorizer()
tf_idf_matrix = model.fit_transform(X)
# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))
```

```
# TF-IDF weighted Word2Vec
```

```

tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in tqdm(list_of_sentence): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
            # tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole corpus
            # sent.count(word) = tf value of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_sent_vectors.append(sent_vec)
    row += 1

```

100%|██████████| 30000/30000 [10:41<00:00, 46.77it/s]

```
tfidf_w2v_vec_google = np.array(tfidf_sent_vectors)
```

```

tfidf_w2v_vecs_norm = preprocessing.normalize(tfidf_w2v_vec_google)
X_train, X_test, y_train, y_test = train_test_split(tfidf_w2v_vecs_norm,y,test_size=0.2,shuffle=False)
X_train, X_cv, y_train, y_cv = train_test_split(X_train, y_train, test_size=0.2)
print("Train Data",X_train.shape)
print("Test Data",X_test.shape)
print("CV Data",X_cv.shape)

```

Train Data (19200, 50)  
Test Data (6000, 50)  
CV Data (4800, 50)

```

from sklearn.svm import SVC
from sklearn.calibration import CalibratedClassifierCV
from sklearn.metrics import roc_auc_score
import math

```

```
C = [10000,1000,100,10,1,0.1,0.01,0.001,0.0001]
```

```

train_auc = []
cv_auc = []

```

```

for i in C:
    model = SVC(C=i)
    clf = CalibratedClassifierCV(model, cv=3)
    clf.fit(X_train,y_train)
    prob_cv = clf.predict_proba(X_cv)[:,-1]
    cv_auc.append(roc_auc_score(y_cv,prob_cv))
    prob_train = clf.predict_proba(X_train)[:,-1]
    train_auc.append(roc_auc_score(y_train,prob_train))
optimal_C= C[cv_auc.index(max(cv_auc))]
C=[math.log(x) for x in C]

```

```

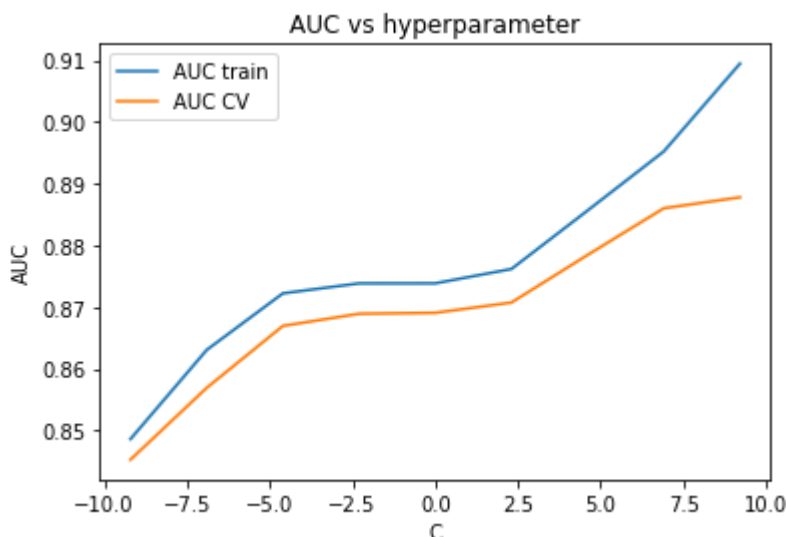
#plot auc vs alpha
x = plt.subplot( )
x.plot(C, train_auc, label='AUC train')
x.plot(C, cv_auc, label='AUC CV')

```



```
plt.title('AUC vs hyperparameter')
plt.xlabel('C')
plt.ylabel('AUC')
x.legend()
plt.show()

print('optimal C for which auc is maximum : ',optimal_C)
```



optimal C for which auc is maximum : 10000

```
from sklearn.metrics import confusion_matrix

#Testing AUC on Test data
model =SVC(C = optimal_C)
clf = CalibratedClassifierCV(model, cv=3)
clf.fit(X_train,y_train)
pred_test = clf.predict_proba(X_test)[:,:1]
fpr1, tpr1, thresholds1 = metrics.roc_curve(y_test, pred_test)
pred_train = clf.predict_proba(X_train)[:,:1]
fpr2,tpr2,thresholds2 = metrics.roc_curve(y_train,pred_train)


#plot ROC curve
x = plt.subplot( )
x.plot(fpr1, tpr1, label = 'Test ROC ,auc='+str(roc_auc_score(y_test,pred_test)))
x.plot(fpr2, tpr2, label='Train ROC ,auc='+str(roc_auc_score(y_train,pred_train)))
plt.title('ROC curve')
plt.xlabel('FPR')
plt.ylabel('TPR')
x.legend()
plt.show()

print("AUC on Test data is " +str(roc_auc_score(y_test,pred_test)))
print("AUC on Train data is " +str(roc_auc_score(y_train,pred_train)))

print("-----")

# Code for drawing seaborn heatmaps
class_names = ['negative','positive']
df_heatmap = pd.DataFrame(confusion_matrix(y_test, pred_test.round()), index=class_names, columns=cl
fig = plt.figure( )
heatmap = sns.heatmap(df_heatmap, annot=True, fmt="d")
```





AUC on Test data is 0.8892091924425987  
AUC on Train data is 0.909460108011407



```
new = [".tfidf W2V", "RBF", 10000, 0.9094, 0.8892].
results_2.loc[3] = new
```

▼ Performance Table

results

↗	Featuraization		Classifier	alpha	Train-AUC	Test-AUC
0	BOW		SGDClassifier-hinge loss	0.0001	0.9624	0.9455
1	tf-idf		SGDClassifier-hinge loss	0.0001	0.9841	0.9586
2	AVG W2V		SGDClassifier-hinge loss	0.0001	0.9034	0.9050
3	tf-idf W2V		SGDClassifier-hinge loss	0.0001	0.8801	0.8785

results\_2

→	Featuraization	Classifier	C	Train-AUC	Test-AUC
0	Bow	RBF	10000	0.9256	0.9100
1	tf-idf	RBF	10000	0.9253	0.9078
2	AVG W2V	RBF	1000	0.9128	0.9058
3	tfidf W2V	RBF	10000	0.9094	0.8892

## ▼ [6] Conclusions

- 1.Support Vector Machine(SVM) gave the best result
- 2.Linear SVM gave best result better then RBF SVM
- 3.Linear SVM gave best result AUC = 0.9586 with tfidf featuraization and alpha = 0.0001
- 4.RBF SVM training time complexity is high compared to Linear SVM